



# PERFORMANCE

# Speedup

- Number of cores =  $p$
- Serial run-time =  $T_{\text{serial}}$
- Parallel run-time =  $T_{\text{parallel}}$



linear speedup

$$T_{\text{parallel}} = T_{\text{serial}} / p$$

# Speedup of a parallel program


$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

# Efficiency of a parallel program

$$E = \frac{S}{p} = \frac{\left( \frac{T_{\text{serial}}}{T_{\text{parallel}}} \right)}{p} = \frac{T_{\text{serial}}}{p \cdot T_{\text{parallel}}}$$

# Amdahl's Law

- Unless virtually all of a serial program is parallelized, the possible speedup is going to be very limited — regardless of the number of cores available.



$$\text{SpeedUp} = \frac{\text{OriginalTime}}{\text{ImprovedTime}} = \frac{1}{\text{ser} + \frac{\text{par}}{\# \text{cores}}} = \frac{1}{\text{ser} + \frac{(1 - \text{ser})}{\# \text{cores}}}$$

# Scalability

- In general, a problem is *scalable* if it can handle ever increasing problem sizes.
- If we increase the number of processes/threads and keep the efficiency fixed without increasing problem size, the problem is *strongly scalable*.
- If we keep the efficiency fixed by increasing the problem size at the same rate as we increase the number of processes/threads, the problem is *weakly scalable*.

# Some Linux commands


- Information about the hardware
  - `cat /proc/cpuinfo`
- Display Linux tasks (processes)
  - `top`
- We're are going to use GCC
  - `gcc --version`



# INTRODUCTION TO PTHREADS



# A Shared Memory System



# Processes and Threads


- A process is an instance of a running (or suspended) program.
- Threads are analogous to a “light-weight” process.
- In a shared memory program a single process may have multiple threads of control.

# POSIX® Threads

- Also known as Pthreads.
- A standard for Unix-like operating systems.
- A library that can be linked with C programs.
- Specifies an application programming interface (API) for multi-threaded programming.

# Caveat

- The Pthreads API is only available on POSIX systems — Linux, MacOS X, Solaris, HPUX, ...



# Hello World! (1)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>


/* Global variable: accessible to all threads */
int thread_count;

void *Hello(void* rank); /* Thread function */

int main(int argc, char* argv[]) {
    long          thread; /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;

    /* Get number of threads from command line */
    thread_count = strtol(argv[1], NULL, 10);

    thread_handles = malloc (thread_count*sizeof(pthread_t));
```



declares the various Pthreads functions, constants, types, etc.

# Hello World! (2)

```
for (thread = 0; thread < thread_count; thread++)  
    pthread_create(&thread_handles[thread], NULL,  
        Hello, (void*) thread);  
  
printf("Hello from the main thread\n");  
  
for (thread = 0; thread < thread_count; thread++)  
    pthread_join(thread_handles[thread], NULL);  
  
free(thread_handles);  
return 0;  
} /* main */
```

# Hello World! (3)

```
void *Hello(void* rank) {  
    long my_rank = (long) rank;  /* Use long in case of 64-bit system */  
  
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);  
  
    return NULL;  
} /* Hello */
```

# Compiling a Pthread program

```
gcc -g -Wall -o pth_hello pth_hello . c -lpthread
```

link in the Pthreads library





# Running a Pthreads program

```
. / pthread_hello <number of threads>
```

```
. / pthread_hello 1
```

Hello from the main thread

Hello from thread 0 of 1

```
. / pthread_hello 4
```

Hello from the main thread

Hello from thread 0 of 4

Hello from thread 1 of 4

Hello from thread 2 of 4

Hello from thread 3 of 4

# Global variables

- Can introduce subtle and confusing bugs!
- Limit use of global variables to situations in which they're really needed.
  - Shared variables.



# Starting the Threads

pthread.h

**One object for  
each thread.**

pthread\_t

```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

# pthread\_t objects

- **Opaque**
- The actual data that they store is system-specific.
- Their data members aren't directly accessible to user code.
- However, the Pthreads standard guarantees that a pthread\_t object does store enough information to uniquely identify the thread with which it's associated.

# A closer look (1)

```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

We won't be using, so we just pass NULL.

Allocate before calling.

# A closer look (2)

```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```


Pointer to the argument that should  
be passed to the function *start\_routine*.

The function that the thread is to run.

# Function started by pthread\_create

- Prototype:  
`void* thread_function ( void* args_p ) ;`
- Void\* can be cast to any pointer type in C.
- So args\_p can point to a list containing one or more values needed by thread\_function.
- Similarly, the return value of thread\_function can point to a list of one or more values.

# Running the Threads



Main thread forks and joins two threads.



# Stopping the Threads

- We call the function `pthread_join` once for each thread.
- A single call to `pthread_join` will wait for the thread associated with the `pthread_t` object to complete.

# Input and Output

- In shared memory programs, only the master thread or thread 0 will access *stdin*.
- In shared memory programs, all the processes/threads can access *stdout* and *stderr*.

# Input and Output

- However, because of the indeterminacy of the order of output to *stdout*, in most cases only a single process/thread will be used for all output to *stdout* other than debugging output.
- Debug output should always include the rank or id of the process/thread that's generating the output.

# Input and Output

- Only a single process/thread will attempt to access any single file other than *stdin*, *stdout*, or *stderr*. So, for example, each process/thread can open its own, private file for reading or writing, but no two processes/threads will open the same file.